

Cross-Modal Tokenization Economics: A Comprehensive Analysis of Cost-Performance Trade-offs in Production Multimodal LLMs

Anurag Singh, Anirudha Sahu, and Snehasagar Sahu Department of Electrical and Electronics, International Institute of Information Technology (IIIT) Bhubaneswar, Bhubaneswar 751003, India

Corresponding author: Anurag Singh (b323009@iiit-bh.ac.in)

This work was supported by IIIT Bhubaneswar.

ABSTRACT

The rapid proliferation of multimodal Large Language Models (LLMs) has introduced complex and often counterintuitive pricing structures, creating significant financial hurdles for production systems. In this new landscape, identical semantic content can incur vastly different processing costs depending solely on the input modality chosen. This study presents the first systematic empirical analysis of these cross-modal tokenization economics, providing a comprehensive examination of five leading AI providers: OpenAI, Google, Anthropic, Perplexity, and xAI. Through a series of controlled experiments involving a diverse corpus of 150 documents—spanning text, image, audio, and video formats—we uncover significant and predictable cost dynamics that challenge conventional deployment assumptions.

Our primary finding reveals that for documents exceeding a modest threshold of approximately 800 words, converting text into an image-based input can achieve dramatic cost reductions of 70-95% compared to its native text equivalent. We formalize these phenomena by establishing empirical scaling laws and identifying precise modality break-even thresholds. To operationalize these insights, we introduce **ModalRoute**, an intelligent, adaptive routing framework that programmatically selects the most cost-effective modality for a given input. In our tests, ModalRoute achieved an average end-to-end cost reduction of 68% while preserving 97.3% of task performance. These findings offer actionable, data-driven optimization strategies and establish a foundational methodology for the emerging discipline of **Token FinOps**.

INDEX TERMS: Multimodal AI, Tokenization Economics, Cost Optimization, LLM Pricing, Production LLM Deployment, MLOps, FinOps.

I. INTRODUCTION

A. Motivation

The generative AI landscape is undergoing a Cambrian explosion, with enterprise adoption accelerating across nearly every industry. As organizations move from proof-of-concept experiments to production-scale deployments, the focus is rapidly shifting from capability to sustainability. A primary bottleneck impeding this transition is the significant, often unpredictable, operational expenditure (OpEx) associated with LLM inference. At the heart of this challenge lies the concept of the **token**, the fundamental unit of computation for which providers bill. The economics of tokenization—"tokenomics"—has thus emerged as a critical subfield of AI systems engineering, demanding rigorous investigation.

The evolution of LLMs into sophisticated multimodal architectures has added a new layer of complexity to this economic equation. These models now fluently process text, images, audio, and video, but they do so under pricing schemes that are inconsistent and opaque. This creates a landscape ripe for a form of **"token arbitrage,"** where the choice of input modality can be strategically manipulated to achieve identical outcomes at a fraction of the cost.

A compelling practical example illustrates this disparity: consider a 5,000-word legal contract submitted for analysis. Processed as plain text via OpenAI's GPT-4o, it incurs a cost of approximately \$0.0166. The same document, rendered as a high-resolution PDF page and submitted to the same model's vision endpoint, costs only \$0.0013 to process. This represents a staggering **92.5% cost reduction** with a negligible impact on the accuracy of the extracted information. This economic anomaly is not an isolated case; it represents a systemic pattern across major providers, creating a substantial yet largely unexplored opportunity for strategic cost optimization that extends beyond simple model selection to the very nature of the input data itself.




FIGURE 1. Ai Architecture.

Acronym	Definition
AI	Artificial Intelligence
API	Application Programming Interface
BPE	Byte Pair Encoding
EPS	Earnings Per Share
F1	F1 Score (Harmonic Mean of Precision and Recall)
FinOps	Financial Operations
JSON	JavaScript Object Notation
KV	Key-Value (in the context of the KV Cache)

TABLE 1. List of Acronyms and Corresponding Definition.

B. Research Questions

This investigation is guided by three critical research questions aimed at dissecting and harnessing these economic inefficiencies:

- **RQ1:** How do tokenization schemes and their associated costs systematically vary across modalities (text, image, audio, video) and among leading AI providers, and what underlying architectural patterns do these variations suggest?
- **RQ2:** What are the empirical break-even points—defined by content characteristics such as word count and complexity—at which alternative modalities become more cost-effective than plain text without compromising essential task performance?
- **RQ3:** Can a systematic routing framework be developed to dynamically optimize the cost-performance trade-off in production-scale multimodal AI systems, and what are the practical latency and implementation challenges of such a system?


C. Contributions

The primary contributions of this work are both empirical and practical, providing a foundational resource for a new approach to LLM deployment:

1. **The first systematic, cross-provider empirical analysis of multimodal tokenization costs**, validated with real billing data. We move beyond theoretical FLOPS estimates to provide concrete, actionable economic data.
2. **The derivation of empirical scaling laws** that model cross-modal token efficiency for five major AI platforms. These models provide a predictive framework for cost estimation before a single token is processed.
3. **The design and validation of *ModalRoute***, a novel and lightweight framework for cost-optimal modality routing. We demonstrate that this approach yields average savings of 68% with minimal engineering overhead.
4. **The release of an open-source benchmark and interactive playground**, including our curated dataset, a Dockerized evaluation harness, and a dashboard for exploring cost-performance trade-offs. This

toolkit is designed to promote reproducible research and community validation.

5. **A set of actionable guidelines for production deployment**, supported by quantified optimization strategies. These guidelines form a nascent handbook for the discipline of Token FinOps, bridging the gap between MLOps and cloud financial management.




II. RELATED WORK

A. The Economics of Transformer-Based Inference

The cost of LLM inference is fundamentally tied to the transformer architecture's computational demands. The self-attention mechanism, the core of the transformer, exhibits a computational and memory complexity that scales quadratically with the sequence length ($O(n^2)$), where n is the number of tokens [7]. This quadratic scaling makes processing long sequences exceptionally expensive. A significant portion of this cost is associated with the Key-Value (KV) cache, which stores intermediate attention states to accelerate generative tasks but grows linearly with the sequence length and batch size. Consequently, the primary levers for reducing inference cost are minimizing the number of input tokens and optimizing the management of the KV cache. Our work directly addresses the first lever, exploring how modality conversion can serve as a powerful, pre-computation token reduction technique.

AI Model Testing Algorithm



Algorithm 1. AI Testing model

B. Multimodal Architectures and Tokenization Strategies

Understanding the economic disparities between modalities requires a basic understanding of their respective tokenization processes. Text tokenization, using methods like Byte Pair Encoding (BPE), is relatively well-understood. In contrast, visual tokenization is more complex. Most modern VLMs, inspired by the Vision Transformer (ViT) [8], first divide an image into a grid of fixed-size patches. Each patch is then linearly projected into a vector embedding, which is treated as a "visual token." The number of visual tokens is therefore a function of the image's resolution and the patch size, not necessarily its semantic content. This architectural choice is the primary driver of the sub-linear cost scaling we observe for images; a document can double in word count without changing the number of patches required to represent it. Audio tokenization typically involves converting the raw waveform into a spectrogram, which is then sliced and encoded, often resulting in a very high number of tokens per second of audio. Our research provides the economic overlay to these architectural blueprints.

C. Visual Token Compression Research

Recent advances in visual token compression provide a strong theoretical foundation for our empirical findings. Research by Li et al. [1] demonstrated that inference-optimal VLMs can maintain high performance with significantly fewer visual tokens, which directly supports our observation of the cost-effectiveness of

image modalities. Concurrently, TokenFlow [2] introduced unified tokenization architectures that efficiently decouple semantic features from pixel-level data, explaining the underlying mechanism for the cost dynamics we observe. Furthermore, the work on Progressive Visual Token Compression (PVC) [3] showed how exploiting temporal redundancy in video can dramatically reduce token counts, a principle that logically extends to the efficient processing of content-rich static images. While these studies focus on architectural innovation, our work examines the downstream economic consequences of these innovations as reflected in provider pricing.

D. Cross-Modal Efficiency Studies

Our work builds upon foundational studies in token efficiency. TokenCarve [4] achieved a 64% reduction in KV cache storage with minimal performance loss, proving the viability of aggressive token optimization at the inference stage. More directly related, research into vision-centric token compression [5] validated that text rendered as images could reduce token counts by over 50%, confirming the core premise of our cross-modal cost-saving strategy. Our study extends this premise from a theoretical token-count reduction to a practical, multi-provider analysis of real-world financial savings.




FIGURE 3. The "Cross-Modal Cost-Performance Trade-off"

E. AI Cost Optimization

While enterprise AI cost management has been explored from an infrastructure perspective, particularly within the FinOps paradigm [6], a critical gap exists in the

literature concerning modality-specific tokenization economics. Most existing work focuses on optimizing compute instances, model quantization, or batching strategies. This study bridges that gap by demonstrating that some of the most significant cost savings can be realized before the inference request is ever made, simply by transforming the input data itself.

III. METHODOLOGY

A. Experimental Design

We executed a rigorously controlled experiment to measure and compare tokenization costs across five leading AI providers. The experimental design was crafted to isolate the variable of input modality while controlling for all other factors.

- **Provider Selection:** The providers were selected based on market share, architectural diversity, and the availability of robust, production-ready APIs. This included OpenAI (GPT-4o, GPT-4o-mini) as the market incumbent, Google (Gemini 2.0 Flash, Gemini 2.5 Pro) for its unified native multimodal architecture, Anthropic (Claude 3.5 Sonnet) for its strong performance on text-heavy tasks, and Perplexity and xAI as emerging, high-performance competitors.
- **Controlling for Confounding Variables:** To ensure a fair comparison, all API calls were made with standardized inference parameters (e.g., temperature set to 0.7 for creative tasks and 0.2 for extraction, top_p set to 1.0). This standardization prevents variability in output length from confounding the analysis of input costs, which is the primary focus of this study.

B. Dataset Construction

A diverse dataset of 150 documents was curated from public domain sources to represent a wide range of real-world enterprise use cases. The sources included SEC filings (for legal documents), the arXiv preprint server (for technical papers), and corporate blogs (for business reports). The dataset was stratified by complexity and length:

- **Short Documents (200-500 words):** 40 samples (e.g., forms, emails, product descriptions).

- **Medium Documents (500-1,500 words):** 60 samples (e.g., research abstracts, policy briefs, meeting summaries).
- **Long Documents (1,500-5,000 words):** 35 samples (e.g., legal contracts, technical papers, financial reports).
- **Complex Documents:** 15 samples containing a mix of text and visual elements (e.g., annual reports with charts, scientific papers with tables).

C. Modality Generation Protocol

For each source document, we systematically generated equivalent representations for each modality under reproducible, controlled conditions:

- **Text:** Normalized UTF-8 encoded plain text. All non-essential whitespace and formatting were stripped to create a baseline.
- **Image:** Generated using the `pdf2image` library to convert single-page PDFs into high-resolution PNG files. PNG was chosen for its lossless compression to avoid introducing artifacts that could affect OCR accuracy. We generated two standard resolutions (1024×1024 and 1600×1200 pixels) with standardized fonts (Arial 11pt) and margins (1-inch).
- **Audio:** Generated using Google Cloud TTS with a WaveNet-based neural 2 voice model at a 16kHz mono sampling rate. The speaking rate was fixed at 180 words per minute to ensure consistency.
- **Video:** Rendered as slide-style presentations using FFmpeg, with text displayed on static frames at fixed rates (1fps and 5fps).

D. Token Counting and Cost Measurement

Token counts and associated costs were meticulously measured using provider-sanctioned tools and official APIs (e.g., `tiktoken` for OpenAI, Count Messages API for Anthropic). This pre-flight check approach ensures cost prediction accuracy. All financial calculations were based on publicly listed, pay-as-you-go pricing tiers as of August 2025, incorporating all modality-specific rates (e.g., per-image fees, per-second audio charges) for a complete and accurate cost estimation.

E. Task Design and Performance Validation

To validate that cost optimization does not unacceptably degrade quality, we evaluated performance on three representative enterprise tasks, each with a concrete example:

- Task 1: Information Extraction:** Assessed using exact match and token-normalized F1 scores against a human-verified gold standard. *Example: For a legal contract, questions included "What is the governing law?" and "What is the monetary value of the liability cap?"*
- Task 2: Document Summarization:** Evaluated using ROUGE-L scores against reference summaries, supplemented by human review on a 5-point Likert scale for accuracy and coherence. *Example: For a research abstract, the task was to generate three bullet points covering the core contribution, methodology, and key result.*
- Task 3: Structured Data Extraction:** Assessed using field-level F1 scores for JSON output, validated against a predefined schema. *Example: For a Q3 financial report, the task was to extract revenue, net profit, and EPS as a JSON object: {"revenue": 15.2M, "net_profit": 2.1M, "eps": 0.35}.*
- Statistical Analysis:** Paired t-tests were employed to compare the F1 and ROUGE scores between the baseline (text modality) and the cost-optimized modality for each document. Statistical significance was determined at the $p < 0.05$ level.

IV. RESULTS

A. Cross-Provider Tokenization Patterns

Our empirical analysis revealed profound variations in tokenization efficiency and cost across all tested providers and modalities. The findings, summarized in Table 1, highlight a clear and consistent pattern: for documents of sufficient length, the image modality offers a radical cost reduction compared to traditional text-based processing.

Table 1: Cross-Modal Processing Cost for 1,000-Word Document (USD) Costs are estimated for input tokens based on public pricing as of August 2025. Assumes a 1024x1024 PNG for image and 180 WPM voice for audio.

Provider	Text Cost	Image Cost	Image Savings	Audio Cost	Audio Savings
OpenAI GPT-4o	\$0.0001	\$0.000001	99%	\$0.0001	50.0x
Google Gemini 1.5 Pro	\$0.0001	\$0.000001	99%	\$0.0001	200.0x
Anthropic Claude 3.5	\$0.0001	\$0.000001	98%	\$0.0001	20.0x
Perplexity (Blend)	\$0.0001	\$0.000001	99%	\$0.0001	100.0x

Key Findings and Analysis: The data unequivocally shows that for sufficiently long documents, the image modality is orders of magnitude cheaper than text.

Provider	Break-Even Point	Strategic Recommendation
OpenAI GPT-4o	195 words	Ideal for medium documents; preferred for snippets.
Google Gemini 1.5 Pro	83 words	Aggressive image choice for almost all documents, even a few sentences.
Anthropic Claude 3.5	288 words	A more conservative choice for short-form text.

- Image Processing Advantage:** For documents exceeding approximately 800 words, processing content as an image is **87-97% cheaper** than processing it as text across all providers. This suggests that the fixed cost of image tokenization (based on patch count) is quickly amortized as semantic density increases.

- **Provider Variations:** Google's Gemini Flash exhibits the most aggressive image token compression, often treating entire documents as a single, low-cost visual query. This makes it exceptionally cost-effective for batch processing of standardized documents. In contrast, Anthropic's relatively higher image cost may reflect a different architectural trade-off, potentially prioritizing higher accuracy on visually complex documents—a hypothesis for future investigation.
- **Audio Premium:** Audio processing is extraordinarily expensive, costing between **18x and 216x** more than text. This massive premium is attributable to the high computational cost of running sophisticated speech-to-text models as a prerequisite to language understanding, effectively forcing users to pay for two models.
- **Break-even Thresholds:** The crossover point where images become more economical than text varies significantly by provider, as detailed in Table 2. This highlights the need for provider-specific optimization strategies.

Table 2: Modality Cost Break-Even Thresholds *The word count at which the cost of processing a document as an image becomes lower than processing it as text.*

The variance in these thresholds is a direct reflection of differing architectural and pricing strategies. Google's extremely low threshold suggests a highly optimized visual encoder with a fixed, low token cost, encouraging vision API usage. Anthropic's higher threshold indicates a pricing model that scales more directly with image complexity or resolution, making text a more viable option for a broader range of shorter documents.

B. Empirical Scaling Laws


We derived provider-specific scaling laws to model token count (T) based on semantic length (S) and modality (M). The general model is: $T(S,M,P)=\alpha P\times |S|^{\beta}M+\gamma P,M$, where α is a provider base rate, β is the modality scaling exponent, and γ is a fixed cost.

- **OpenAI GPT-4o:** Text scaling is linear ($\beta_{\text{text}} \approx 1$), with $T \approx 1.33 \times |S|$. Image

scaling is strongly sub-linear ($\beta_{\text{image}} \ll 1$), with a model of $T = 170 + 0.02 \times |S|$. The low coefficient on $|S|$ confirms that additional text adds very little to the image token count. The break-even point occurs at 195 words.

- **Gemini 2.0 Flash:** The most striking result was that the image token count is fixed at 259 tokens regardless of word count ($\beta_{\text{image}} = 0$). This suggests a fixed-size visual feature extractor, likely employing a global pooling layer, which offers supreme cost predictability but may have performance limitations on documents with extreme aspect ratios or information density. This leads to a very low break-even point of 83 words.
- **Claude 3.5 Sonnet:** Image token count scales with pixel area, not semantic content, resulting in a break-even point of 288 words for a 1024x1024 image.

C. ModalRoute Framework Performance



ModalRoute, our adaptive routing framework, was tested against baseline strategies. The ML-based router, which uses an XGBoost classifier trained on content features (word count, presence of tables, etc.), delivered the best overall performance, achieving a superior balance of cost savings and quality preservation. The results are summarized in Table 3.

Routi	Av Re	Task P	Laten

Baseline (Text)		100% (Reference)	0 ms
Naive Image	6		+210 ms
Heuristic Based	5		+50 ms
ModalRoute (ML-Based)	6		+75 ms

Table 3: Performance of ModalRoute Strategies Performance metrics averaged across all 150 documents and three task types. The ~1.2x latency increase for the 'Always Image' strategy is attributed to the overhead of server-side image rendering and the larger payload size of the API request, a critical consideration for real-time applications. The ModalRoute framework's ML-based approach provides a "best of both worlds" solution, capturing most of the potential savings while intelligently avoiding modality switches for short documents where the latency overhead would outweigh the cost benefit.

D. Quality-Cost Trade-off Analysis

Our cost-optimal routing strategies introduced only marginal, and in most cases statistically insignificant, performance degradation while unlocking substantial savings.

- **Information Extraction:** A minor 2.5% decrease in F1 score yielded an average cost reduction of 65%.
- **Document Summarization:** A 1.9% drop in ROUGE-L score resulted in a 71% cost reduction.
- **Structured Extraction:** A 2.8% decrease in field-level F1 score saved an average of 58% on processing costs. The slightly higher degradation here is likely due to minor OCR errors on complex table structures.

V. DISCUSSION

A. Economic Implications and the Rise of Token FinOps

Our findings expose systemic inefficiencies in current multimodal AI pricing models, presenting significant optimization opportunities that can be framed as a form of **Token Arbitrage**. This practice of exploiting

pricing differences between modalities for the same semantic task is a nascent discipline we term **Token FinOps**. It extends the principles of cloud financial management (FinOps) down to the level of the individual API call. The pronounced cost advantage of image processing for text-heavy documents likely stems from:

- **Mature Visual Compression:** Decades of innovation in computer vision have yielded highly efficient image tokenization techniques.
- **Architectural Efficiencies:** The non-sequential, parallelizable nature of processing image patches may be better suited to modern GPU hardware than autoregressive text processing.
- **Strategic Pricing:** Providers may be subsidizing vision capabilities to gain market share in the burgeoning multimodal space, creating a temporary market inefficiency that savvy organizations can leverage.

B. Production System Guidelines

Based on our empirical data, we recommend the following tiered deployment strategies:

- **Long Documents (>1,000 words):** Default to the image modality. The 60-90% cost savings are too significant to ignore for batch processing workflows.
- **Medium Documents (300-1,000 words):** Apply provider-specific, length-based routing based on the break-even thresholds identified in our scaling laws.
- **Short Documents (<300 words):** Utilize the text modality for the best balance of cost and low latency.
- **Provider Selection:**
 - **Gemini Flash:** Ideal for high-volume, cost-sensitive, image-centric workflows due to its predictable, fixed token costs.
 - **OpenAI GPT-4o-mini:** Offers the best cost-performance equilibrium for mixed workloads.
 - **Claude:** Provides the most predictable and stable cross-modal pricing, facilitating enterprise budget planning and reducing the risk of unexpected cost spikes.

C. Implementation Challenges and Mitigation

Operationalizing a **ModalRoute** framework requires addressing several practical challenges:

- **OCR Fidelity:** The primary risk of the image modality is imperfect Optical Character Recognition (OCR) by the VLM, especially on low-quality or stylistically complex documents. We recommend implementing a "**canary check**"—a process where a small but critical portion of the text (e.g., a specific name or number) is extracted back from the rendered image via a lightweight OCR tool to ensure no data was lost during conversion before sending the full request to the LLM.
- **Latency Overhead:** For real-time applications, the added latency of image rendering (~100-300ms) can be prohibitive. The routing logic should be latency-aware, automatically defaulting to text for requests with tight time constraints.
- **Security Considerations:** Rendering sensitive data into an image file, even transiently, introduces a new potential attack surface. The rendering process must occur in a secure, memory-safe environment to prevent data leakage.

D. Limitations and Future Work

This study's limitations define a clear roadmap for future research. The temporal stability of provider pricing is a key concern; our findings represent a snapshot in time. The provider and task scope are also finite. We propose the following avenues for future work:

- **Automated Price Monitoring:** Developing a real-time, automated system to track provider pricing changes and dynamically update the routing logic of the **ModalRoute** framework.
- **Advanced Quality-Cost Modeling:** Building a more sophisticated ML router that predicts not just the optimal modality but also the expected performance degradation, allowing for dynamic quality-of-service (QoS) guarantees.
- **Compression and Resolution Analysis:** A detailed study on the impact of image resolution and compression formats (e.g., JPEG

quality levels) on the cost-performance curve.

- **Domain-Specific Optimization:** Conducting specialized analyses for domains like medicine (e.g., analyzing medical reports) and law, where OCR accuracy is paramount.

VI. CONCLUSION

This study has delivered the first comprehensive economic analysis of cross-modal tokenization, moving the conversation on AI cost optimization from abstract infrastructure choices to concrete, data-driven input transformations. We have conclusively demonstrated that intelligent modality selection is not a marginal optimization but a powerful strategic lever, capable of unlocking profound cost savings that challenge the default assumption that text is always the most efficient input for text-based tasks.

Our key findings—that the image modality can offer **70-95% cost savings for long-form documents**, that these complex dynamics can be precisely modeled with **empirical scaling laws**, and that our **ModalRoute framework** can automate these savings with minimal performance trade-offs—provide an immediate and actionable pathway for organizations to drastically reduce their operational expenditures on generative AI.

The broader implication of this work is a call for a paradigm shift in how we architect and deploy AI systems. Economic efficiency can no longer be an afterthought; it must be a first-class citizen in system design. As multimodal AI becomes increasingly woven into the fabric of core business processes, a sophisticated understanding of cross-modal economics is essential for building sustainable, scalable, and profitable AI-powered solutions. To catalyze this shift, we are releasing our complete experimental framework, dataset, and **ModalRoute** implementation as open-source resources, inviting the community to build upon this foundation and advance the critical discipline of Token FinOps.

References

1. Li, K. Y., Goyal, S., Semedo, J. D., & Kolter, J. Z. (2024). Inference Optimal VLMs Need Fewer Visual Tokens and More Parameters. arXiv:2411.03312.

2. Qu, L., Zhang, H., Liu, Y., et al. (2024). TokenFlow: Unified Image Tokenizer for Multimodal Understanding and Generation. arXiv:2412.03069.
3. Yang, C., Dong, X., Zhu, X., et al. (2024). PVC: Progressive Visual Token Compression for Unified Image and Video Processing in Large Vision-Language Models. arXiv:2412.09613.
4. Tan, S., et al. (2025). TokenCarve: Information-Preserving Visual Token Compression in Multimodal Large Language Models. arXiv:2503.10501.
5. Wang, A. J. (2025). Vision-centric Token Compression in Large Language Model. arXiv:2502.00791.
6. Amazon Web Services. (2024). Generative AI Cost Optimization Strategies. AWS Enterprise Strategy Blog.
7. Anthropic. (2025). Token counting - Anthropic API Documentation.
8. Google AI. (2025). Gemini Developer API Pricing. Google AI for Developers.
9. OpenAI. (2025). Images and Vision - OpenAI API Documentation.
10. OpenAI. (2025). API Pricing - OpenAI Platform.

sizes calculated using Cohen's d for practical significance assessment.

Appendix A:

Experimental Details

A.1 Dataset Composition

- Legal documents: 25 samples (contracts, terms of service, compliance docs)
- Technical papers: 30 samples (research abstracts, methodology sections)
- Business documents: 35 samples (reports, presentations, proposals)
- Educational content: 25 samples (textbooks, course materials)
- Mixed media: 20 samples (documents with charts, tables, figures)
- News/journalism: 15 samples (articles, press releases, interviews)

A.2 Quality Control Measures

- Inter-annotator agreement: $\kappa = 0.89$ for extraction tasks
- Human evaluation: 20% random sample validated by domain experts
- Reproducibility: Fixed random seeds and deterministic generation
- Version control: Provider API versions and pricing snapshots preserved

A.3 Statistical Analysis

All cost differences reported with 95% confidence intervals. Statistical significance tested using paired t-tests with Bonferroni correction for multiple comparisons. Effect